



FIELD GUIDE 02

The Clinician's AI Prompting Guide

CLINICALLY
INTELLIGENT

FIELD GUIDE 02


AI PROMPTING GUIDE



THE WEEKLY
NEWSLETTER FOR
HEALTHCARE
PROFESSIONALS
WHO REFUSE TO
BE LEFT BEHIND.

CLINICALLYINTELLIGENT.COM

2026



A practical field guide to
writing better AI prompts
for healthcare.

Built for clinicians.
Guided by judgement.

This is a **guide**, not a template library. —

Read the six questions once.
Return only when a prompt isn't
producing **usable output**.

THIS GUIDE RUNS IN THREE STAGES:

01

What every prompt needs

Three checks before you hit enter.

02

What good prompts do

Three moves that bad prompts miss.

03

When to stop and try again

Three signals the prompt is not working.

Each stage contains three questions.
If your prompt is not producing usable output,
one of these six checks is the reason.

WHEN A PROMPT FAILS:



Do not use it.

One of the six checks failed.



Find the failed check.

Fix it. Rerun the prompt.



Still unsure?

Ask your IG Lead before improving.



IMPORTANT



Use alongside the Safety Check.
This guide assumes you've read
the AI Safety Check (Issue 07).



Same framework every time.
All prompting techniques here
follow the same IG framework.



Composite scenarios only.
Never upload identifiable
patient information.



Every output is a draft.
You review. You verify.
You sign off.



Trust systems, not tools.
Real patient documentation
stays in your Trust's systems.
Not in any AI tool.



NOT READ THE SAFETY CHECK?

Stop. Get it at
clinicallyintelligent.com
before using this guide.



What every prompt needs. —

Three things every clinical prompt must contain.
Answer **yes** to all three before you hit enter.

THE BRIEFING TEST

If you were briefing a junior colleague on the same task, what would you tell them?

Tell the AI the same thing.

CHECK 01

Have I named the role and the clinical context?

The AI does not know it is helping a clinician until you tell it.

Naming your role and the setting changes the output register from generic to clinical.

The more specific the role and setting, the more clinically appropriate the output.

GENERIC

"Help me write a document."

CLINICAL

"Act as a clinical documentation assistant helping an NHS Occupational Therapist in a community frailty service."

CHECK 02

Have I given enough context for the output to be useful?

Vague prompts produce vague outputs.
Specific prompts produce specific outputs.

The context should include the audience for the output, the format you need, and the clinical scenario the output will apply to.

A composite scenario is fine and expected.

CONTEXT EXAMPLES

- Patient scenario
- Setting / service
- Audience
- Purpose of the output
- Key constraints

Give enough for clarity. Nothing that breaks IG.

CHECK 03

Have I specified what the output must include and exclude?

Tell the AI what to produce and what to leave out.

READING AGE

SOURCE REQUIREMENTS

STRUCTURE

LENGTH

TONE

CONTENT YOU DO NOT WANT

Constraints improve output quality. The AI cannot guess your preferences. If it matters to you, put it in the prompt.

EXAMPLES

- Use UK English
- Reading age: 12
Use —
- Use NICE guideline NG31
- Exclude diagnostic advice
- Format as bullet list
- Maximum 150 words

Be explicit.
Get better output.

What good prompts do.

Three moves that separate professional prompts from amateur ones.

Answer **yes** to all three before you commit to using the output.

Every prompt should leave the clinical decision with you.

CHECK

04

Am I asking the AI to accelerate my thinking or replace it?

Prompts that ask the AI to make clinical decisions produce confidently wrong output. Prompts that ask the AI to scaffold your thinking produce useful output.



DO

- ✓ Draft
- ✓ Structure
- ✓ Explore options
- ✓ Ask questions
- ✓ Scaffold thinking



DON'T

- ✗ Decide
- ✗ Recommend
- ✗ Diagnose
- ✗ Choose

The clinical judgement stays with you. Every time.

CHECK

05

Am I requiring evidence, sources, or reasoning?

Quality output is built on verifiable input. Always require evidence.

- ✓ Where did this come from?
- ✓ Can I verify it?
- ✓ Does it cite guidance?
- ✓ Is the guidance current?

BEFORE A CLAIM ENTERS CLINICAL WORK:



Require a source in the prompt



Ask the AI to show its working



Verify before you use it

CHECK

06

Have I told the AI what to do if it does not know?

Do not let the AI guess or fill in gaps. Tell it what to do instead.

- Write "unknown" or "not available"
- Ask for what is missing
- Do not fabricate information

The prompt should never leave the AI to fill gaps on its own.

FIELD NOTE

When creating templates, always tell the AI what to do if information is missing. Otherwise, it will fill the gaps itself.

POSITION THE AI CORRECTLY

It is your assistant, not your replacement.

The clinical judgement stays with you.

Every time.

THE GOAL

Use AI



Better prompts



Better thinking



Better care



Better outcomes



CLINICALLY
INTELLIGENT

FIELD GUIDE 02
AI PROMPTING
PLAYBOOK

THE WEEKLY
NEWSLETTER FOR
HEALTHCARE
PROFESSIONALS
WHO REFUSE TO
BE LEFT BEHIND.

CLINICALLYINTELLIGENT.COM

2026

When to stop and **try again.**

Bad outputs are usually prompt problems, not AI problems.
Identify the signal. Improve the prompt. Run it again.
Answer **yes** to all three before moving on.

CI

Good prompting isn't getting the right output the first time. It's improving the prompt until the output earns your trust.

SIGNAL

01

Is the output too generic?

If the output could apply to almost any patient, your prompt wasn't specific enough.

TRY THIS:

- ✓ Add more clinical context
- ✓ Name the audience
- ✓ Specify the format
- ✓ Run it again



If it could apply to everyone, it probably helps no one.

SIGNAL

02

Can I verify every claim?

If you can't trace the source, don't use the answer.

TRY THIS:

- ✓ Require sources
- ✓ Ask for reasoning
- ✓ Check every citation
- ✓ Run it again



No source = No clinical use.

SIGNAL

03

Does something feel off?

Trust your clinical instinct. If something feels wrong, improve the prompt—not your standards.

TRY THIS:

- ✓ Tighten the instructions
- ✓ Add constraints
- ✓ Change the perspective
- ✓ Run it again



Clinical judgement is data. Don't ignore it.

THE ITERATION LOOP

Output not usable?



Identify the signal



Improve the prompt
(change ONE part)



Run again



Usable?



YES

Review • Verify • Use



NO

Repeat the loop

THE ITERATION TOOLKIT

Change one thing at a time.



Context



Task



Constraints



Perspective



Structure



REMEMBER

A poor output is feedback about the prompt—not permission to lower your standards.

**Iterate.
Verify.
Use.**

Same task. Different results.

Three real clinical prompting examples.
Steal the structure, not the words.



Don't copy
the words.
*Copy the
structure.*

01 Example one • Clinical documentation

BAD

"Write a discharge
summary."

WHY THIS WORKS

- ✓ Defines the role
- ✓ Gives clinical context
- ✓ Specifies the output
- ✓ Prevents invention
- ✓ Handles missing information

GOOD

STEAL THIS STRUCTURE

Act as a clinical documentation assistant helping an NHS
Occupational Therapist. Draft a discharge summary for the
following composite scenario. Include assessment findings,
suitable for a receiving community therapist. Do not
invent patient details. If anything is missing write
'to be personalised by clinician.'
Scenario: [composite scenario].'

Defines the role

Sets the
audience

Specifies the
output

Prevents
invention

Handles gaps
safely

STEAL THIS STRUCTURE

ROLE
Define who AI
should be.

CONTEXT
Give the clinical
setting and
audience.

OUTPUT
State exactly
what you want
in the output.

CONSTRAINTS
Tell AI what
not to do and
how to behave.

GAP HANDLING
Tell AI what to
do if something
is missing.

EVIDENCE
Ask for sources
and how current
they must be.

Use every time.

02 Example two • Research and evidence

BAD

"What is the latest
guidance on falls
prevention."

WHY THIS WORKS

- ✓ Names the population and geography
- ✓ Requires guideline number and date
- ✓ Requires verifiable sources
- ✓ Flags outdated guidance(s)
- ✓ Structures for clinical use

GOOD

STEAL THIS STRUCTURE

Summarise the current UK NICE guidance on falls
prevention in community-dwelling older adults.
Include the specific guideline number and date of
last update. Cite sources I can verify with
direct links. Flag if any guidance is more than
two years old. Structure as: overview, key
recommendations, red flags, escalation criteria.'

Defines
the task

Sets the
population
and geography

Requires
evidence

Flags outdated
guidance(s)

Structures for
clinical use



Don't memorise prompts. Memorise the pattern.

Thirty extra seconds writing the prompt saves minutes reviewing the output.

Use both.

SAFETY CHECK
Is it safe?
Check for hallucinations,
bias, and patient safety
risks first.



PROMPT GUIDE
Is it useful?
Use the structure
to get clear, high-quality,
usable outputs.

FIELD GUIDE 01

**The Clinician's
AI Safety Check.**

Read before using AI clinically.



Free at
clinicallyintelligent.com

Run both. Safety first, **quality second**, every time.

Example three • Patient-facing material.

The same task.

A vague prompt gets
a generic result.

A structured prompt
gets something
you can use.



BAD

"Write an exercise
programme for a patient."



GOOD

STEAL THIS STRUCTURE


'Act as a **clinical design assistant** helping an OT' → Defines the role
produce a home exercise programme for a → Sets the reading
composite patient. **Reading age 11.** Plain English.
Warm tone speaking directly to the patient and → Controls the tone
the family member supporting them. Use only
the exercises I specify below. If any detail is → Specifies the
missing write 'to be personalised by clinician' in
that field. **Do not invent patient details** or add → Prevents
exercises I have not listed. **Exercises: [list with**
reps, sets, frequency].' → Handles missing
information



Every good prompt looks different.
Every great prompt follows the same structure.

The pattern.

Six moves that turn
a vague request into
a useful result.



Don't memorise prompts. Memorise the pattern.

Thirty extra seconds writing the prompt saves minutes reviewing the output.

Use both.



SAFETY CHECK
Is it safe?

Check for hallucinations, bias, and patient safety risks first.



PROMPT GUIDE
Is it useful?

Use the structure to get clear, high-quality, usable outputs.

Run both. Safety first, **quality second**, every time.

FIELD GUIDE 01

**The Clinician's
AI Safety Check.**

Read before using AI clinically.



Free at
clinicallyintelligent.com

The technology is new. Your responsibilities aren't.

Every prompt in this guide assumes the same **clinical responsibilities** you already work under.



01

Use composite scenarios only.

No identifiable patient information enters consumer AI at any point.



02

Every output is a draft.

You review. You verify.
You sign off.

The AI is not signing anything. You are.



03

Verify every source.

AI can invent citations or link to outdated guidance.
Verification is **your** responsibility.



04

Keep patient records in Trust systems.

Real patient documentation belongs in your Trust's systems.

Not in any consumer AI tool.



CONSUMER AI

Not assessed against NHS clinical safety standards.

✗ DCB0129

✗ DCB0160

Treat anything you upload as potentially processed outside the standard NHS health IT is normally held to.



If you are unsure whether a specific prompt or use case sits within the framework, ask your IG lead **before** you run it.



Everything you've learned in this guide assumes these four principles.
They are not optional.



BEFORE USING AI CLINICALLY

Run the Safety Check.

- ✓ Information governance
- ✓ Clinical safety
- ✓ Risk assessment



Read this
before you use AI.

clinicallyintelligent.com



✓ Good prompting improves quality. Good governance protects patients. **You need both.**

Better prompts. Better care.

Good prompts protect quality.
Clinical judgement protects patients.

AI will keep changing.
Clinical judgement remains
your responsibility.

The best clinicians won't be
those who use AI the most.

They'll be the ones who know
when to trust it—
and when not to.

AI can generate words.
Only clinicians deliver care.



BEFORE EVERY
CLINICAL PROMPT

**Run the
Safety Check.**

A quick check.
A safer use of AI.



clinicallyintelligent.com



Built for clinicians. Guided by judgement.

YOUR RESPONSIBILITIES

Your responsibilities don't change
because the technology does.



Review every output

Check accuracy and relevance.



Apply your own clinical judgement

You know your patient. You make
the decision.



Follow your organisation's policies and information governance requirements

Keep data secure and compliant.



You remain responsible for every clinical decision.

Better prompts. Better care.

Good prompts protect quality.
Clinical judgement protects patients.

AI can generate words.
Only clinicians deliver care.

Thank you.

*To every clinician choosing to use AI thoughtfully,
carefully, and in the interests of better patient care.*



CLINICALLY
INTELLIGENT



Written & designed by
Clinically Intelligent



Design
CI Design Studio



Version
1.0 • July 2026



Contact
hello@clinicallyintelligent.com